

CT Scan Image Classification

▼ CT Scan Image Classification

▼ Connect to Google Drive to access Dataset

```
[ ] from google.colab import drive
drive.mount('/content/drive')
%cd '/content/drive/My Drive/Colab Notebooks/Data/'
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).
/content/drive/My Drive/Colab Notebooks/Data

▼ Import Libraries

```
from builtins import range, input

from tensorflow.keras.layers import Input, Lambda, Dense, Flatten, AveragePooling2D, Dropout
from tensorflow.keras.models import Model, load_model
from tensorflow.keras.applications.resnet50 import ResNet50
from tensorflow.keras.applications.resnet50 import preprocess_input
from tensorflow.keras.preprocessing import image
from tensorflow.keras.preprocessing.image import ImageDataGenerator

from sklearn.metrics import confusion_matrix, roc_curve
import seaborn as sns
import numpy as np
import matplotlib.pyplot as plt

import cv2
from glob import glob

from keras.callbacks import ModelCheckpoint, EarlyStopping
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelBinarizer
from tensorflow.keras.utils import to_categorical
```

```
#define size to which images are to be resized
IMAGE_SIZE = [80, 80]

#define paths
covid_path = '/content/drive/MyDrive/Colab Notebooks/Data/COVID'
noncovid_path = '/content/drive/MyDrive/Colab Notebooks/Data/non-COVID'

# Use glob to grab images from path .jpg or jpeg
covid_files = glob(covid_path + '/*')
noncovid_files = glob(noncovid_path + '/*')
```

▼ Resize images

```
[ ] covid_labels = []
    noncovid_labels = []

    covid_images=[]
    noncovid_images=[]

    for i in range(len(covid_files)):
        image = cv2.imread(covid_files[i])           # read file
        image = cv2.resize(image,(80,80))           # resize as per model for covid
        covid_images.append(image)                   # append image
        covid_labels.append('COVID')                 #append class label

    for i in range(len(noncovid_files)):
        image = cv2.imread(noncovid_files[i])
        image = cv2.resize(image,(80,80))
        noncovid_images.append(image)
        noncovid_labels.append('non-COVID')
```

▼ Visualize First 25 images from dataset


```
▶ # look at a random image for fun
def plot_images(images, title):
    nrows, ncols = 5, 5
    figsize = [6, 6]

    fig, ax = plt.subplots(nrows=nrows, ncols=ncols, figsize=figsize, facecolor=(1, 1, 1))


    for i, axi in enumerate(ax.flat):
        axi.imshow(images[i])
        axi.set_axis_off()

    plt.suptitle(title, fontsize=24)
    plt.tight_layout(pad=0.2, rect=[0, 0, 1, 0.9])
    plt.show()
plot_images(covid_images, 'COVID-19 Positive CT Scan')
plot_images(noncovid_images, 'COVID-19 Negative CT Scan')
```

COVID-19 Positive CT Scan



COVID-19 Negative CT Scan



▼ Normalization

```
[ ] # Convert to array and Normalize to interval of [0,1]
    covid_images = np.array(covid_images) / 255
    noncovid_images = np.array(noncovid_images) / 255
```

▼ Train Test Splitting

```
▶ # Split into training and testing sets for both types of images

covid_x_train, covid_x_test, covid_y_train, covid_y_test = train_test_split(
    covid_images, covid_labels, test_size=0.2, random_state=1)

noncovid_x_train, noncovid_x_test, noncovid_y_train, noncovid_y_test = train_test_split(
    noncovid_images, noncovid_labels, test_size=0.2, random_state=1)


# Merge sets for both types of images
X_train = np.concatenate((noncovid_x_train, covid_x_train), axis=0)
X_test = np.concatenate((noncovid_x_test, covid_x_test), axis=0)
y_train = np.concatenate((noncovid_y_train, covid_y_train), axis=0)
y_test = np.concatenate((noncovid_y_test, covid_y_test), axis=0)


# Make labels into categories - either 0 or 1, for our model
y_train = LabelBinarizer().fit_transform(y_train)
y_train = to_categorical(y_train)
y_test = LabelBinarizer().fit_transform(y_test)
y_test = to_categorical(y_test)

# y_train and y_test contain class labels 0 and 1 representing COVID and NonCOVID for X_train and X_test
```


```
▶ plot_images(covid_x_train, 'X_train')
plot_images(covid_x_test, 'X_test')
```

X_train





X_test



▼ Model Building

```
[ ] # Building Model
    from tensorflow import keras
    resnet = ResNet50(weights="imagenet", include_top=False,
        input_tensor=Input(shape=(80, 80, 3)))

    outputs = resnet.output
    outputs = Flatten(name="flatten")(outputs)
    outputs = Dropout(0.3)(outputs)
    outputs = Dense(2, activation="sigmoid")(outputs)


    model = Model(inputs=resnet.input, outputs=outputs)

    for layer in resnet.layers:
        layer.trainable = False

    model.compile(
        loss='binary_crossentropy',
        optimizer=keras.optimizers.Adamax(learning_rate=0.001),
        metrics=['accuracy']
    )
```

```
# Define callbacks
early_stopping = EarlyStopping(
    monitor='val_loss',
    patience=5,
    verbose=1,
    restore_best_weights=True
)

model_checkpoint = ModelCheckpoint(
    'model_checkpoint.h5',
    monitor='val_loss',
    save_best_only=True,
    verbose=1
)
```

 `model.summary()`

▼ Data Augmentation

```
[ ] datagen = ImageDataGenerator(rotation_range=20,
    width_shift_range=0.2,
    height_shift_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True
)
```


▼ Training the Model

```
[ ] batch_size=2
    epochs=30

hist = model.fit(datagen.flow(X_train, y_train, batch_size=batch_size),
                  validation_data=(X_test, y_test),
                  validation_steps=len(X_test) / batch_size,
                  steps_per_epoch=len(X_train) / batch_size,
                  epochs=epochs, callbacks=[early_stopping, model_checkpoint]
                )
```

```
Epoch 1/30
991/993 [=====>.] - ETA: 0s - loss: 0.6901 - accuracy: 0.6044
Epoch 1: val_loss improved from inf to 0.59253, saving model to model_checkpoint.h5
993/993 [=====] - 23s 20ms/step - loss: 0.6897 - accuracy: 0.6047 - val_loss: 0.5925 - val_accuracy: 0.6419
Epoch 2/30
990/993 [=====>.] - ETA: 0s - loss: 0.6860 - accuracy: 0.6116
Epoch 2: val_loss did not improve from 0.59253
993/993 [=====] - 17s 17ms/step - loss: 0.6853 - accuracy: 0.6123 - val_loss: 0.6269 - val_accuracy: 0.6519
Epoch 3/30
991/993 [=====>.] - ETA: 0s - loss: 0.6722 - accuracy: 0.6302
Epoch 3: val_loss did not improve from 0.59253
993/993 [=====] - 16s 16ms/step - loss: 0.6715 - accuracy: 0.6304 - val_loss: 0.8262 - val_accuracy: 0.5855
Epoch 4/30
992/993 [=====>.] - ETA: 0s - loss: 0.6733 - accuracy: 0.6280
Epoch 4: val_loss did not improve from 0.59253
993/993 [=====] - 16s 16ms/step - loss: 0.6733 - accuracy: 0.6284 - val_loss: 0.8341 - val_accuracy: 0.5775
Epoch 5/30
991/993 [=====>.] - ETA: 0s - loss: 0.6508 - accuracy: 0.6433
Epoch 5: val_loss improved from 0.59253 to 0.57693, saving model to model_checkpoint.h5
993/993 [=====] - 18s 18ms/step - loss: 0.6503 - accuracy: 0.6440 - val_loss: 0.5769 - val_accuracy: 0.6962
Epoch 6/30
993/993 [=====] - ETA: 0s - loss: 0.6611 - accuracy: 0.6339
Epoch 6: val_loss did not improve from 0.57693
993/993 [=====] - 17s 17ms/step - loss: 0.6611 - accuracy: 0.6339 - val_loss: 0.6241 - val_accuracy: 0.6559
Epoch 7/30
992/993 [=====>.] - ETA: 0s - loss: 0.6523 - accuracy: 0.6361
Epoch 7: val_loss did not improve from 0.57693
993/993 [=====] - 16s 16ms/step - loss: 0.6520 - accuracy: 0.6365 - val_loss: 0.6993 - val_accuracy: 0.6217
Epoch 8/30
989/993 [=====>.] - ETA: 0s - loss: 0.6457 - accuracy: 0.6481
Epoch 8: val_loss improved from 0.57693 to 0.55151, saving model to model_checkpoint.h5
993/993 [=====] - 17s 18ms/step - loss: 0.6457 - accuracy: 0.6485 - val_loss: 0.5515 - val_accuracy: 0.7284
993/993 [=====] - 17s 18ms/step - loss: 0.6457 - accuracy: 0.6485 - val_loss: 0.5515 - val_accuracy: 0.7284
Epoch 9/30
990/993 [=====>.] - ETA: 0s - loss: 0.6566 - accuracy: 0.6364
Epoch 9: val_loss did not improve from 0.55151
993/993 [=====] - 16s 16ms/step - loss: 0.6563 - accuracy: 0.6370 - val_loss: 0.5785 - val_accuracy: 0.6861
Epoch 10/30
993/993 [=====] - ETA: 0s - loss: 0.6366 - accuracy: 0.6475
Epoch 10: val_loss did not improve from 0.55151
993/993 [=====] - 17s 17ms/step - loss: 0.6366 - accuracy: 0.6475 - val_loss: 0.6372 - val_accuracy: 0.6499
Epoch 11/30
993/993 [=====] - ETA: 0s - loss: 0.6583 - accuracy: 0.6410
Epoch 11: val_loss improved from 0.55151 to 0.54613, saving model to model_checkpoint.h5
993/993 [=====] - 18s 18ms/step - loss: 0.6583 - accuracy: 0.6410 - val_loss: 0.5461 - val_accuracy: 0.7364
Epoch 12/30
991/993 [=====>.] - ETA: 0s - loss: 0.6523 - accuracy: 0.6398
Epoch 12: val_loss did not improve from 0.54613
993/993 [=====] - 16s 16ms/step - loss: 0.6526 - accuracy: 0.6400 - val_loss: 0.6415 - val_accuracy: 0.6519
Epoch 13/30
990/993 [=====>.] - ETA: 0s - loss: 0.6521 - accuracy: 0.6323
Epoch 13: val_loss did not improve from 0.54613
993/993 [=====] - 16s 16ms/step - loss: 0.6515 - accuracy: 0.6334 - val_loss: 0.5796 - val_accuracy: 0.6740
Epoch 14/30
991/993 [=====>.] - ETA: 0s - loss: 0.6479 - accuracy: 0.6438
Epoch 14: val_loss did not improve from 0.54613
993/993 [=====] - 17s 17ms/step - loss: 0.6472 - accuracy: 0.6445 - val_loss: 0.7726 - val_accuracy: 0.5956
Epoch 15/30
991/993 [=====>.] - ETA: 0s - loss: 0.6230 - accuracy: 0.6680
Epoch 15: val_loss improved from 0.54613 to 0.53028, saving model to model_checkpoint.h5
993/993 [=====] - 18s 18ms/step - loss: 0.6231 - accuracy: 0.6682 - val_loss: 0.5303 - val_accuracy: 0.7404
```

```

Epoch 16/30
991/993 [=====>.] - ETA: 0s - loss: 0.6438 - accuracy: 0.6448
Epoch 16: val_loss did not improve from 0.53028
993/993 [=====] - 16s 16ms/step - loss: 0.6436 - accuracy: 0.6450 - val_loss: 0.5368 - val_accuracy: 0.7203
Epoch 17/30
992/993 [=====>.] - ETA: 0s - loss: 0.6497 - accuracy: 0.6477
Epoch 17: val_loss did not improve from 0.53028
993/993 [=====] - 16s 16ms/step - loss: 0.6493 - accuracy: 0.6480 - val_loss: 0.7487 - val_accuracy: 0.6076
Epoch 18/30
990/993 [=====>.] - ETA: 0s - loss: 0.6360 - accuracy: 0.6535
Epoch 18: val_loss did not improve from 0.53028
993/993 [=====] - 16s 16ms/step - loss: 0.6361 - accuracy: 0.6531 - val_loss: 0.6345 - val_accuracy: 0.6479
Epoch 19/30
990/993 [=====>.] - ETA: 0s - loss: 0.6343 - accuracy: 0.6475
Epoch 19: val_loss did not improve from 0.53028
993/993 [=====] - 16s 16ms/step - loss: 0.6354 - accuracy: 0.6470 - val_loss: 0.8400 - val_accuracy: 0.5674
Epoch 20/30
993/993 [=====] - ETA: 0s - loss: 0.6493 - accuracy: 0.6490Restoring model weights from the end of the best epoch: 15.

Epoch 20: val_loss did not improve from 0.53028
993/993 [=====] - 16s 16ms/step - loss: 0.6493 - accuracy: 0.6490 - val_loss: 0.7019 - val_accuracy: 0.6258
Epoch 20: early stopping

```

```

▶ # Save Model and Weights
model.save('resnet_ct.h5')
model.save_weights('resnet_weights_ct.hdf5')

```

```

[ ] # Load saved model
model = load_model('resnet_ct.h5')
final_loss, final_accuracy = model.evaluate(X_test, y_test)
print('Final Loss: {}, Final Accuracy: {}'.format(final_loss, final_accuracy))

```

```

16/16 [=====] - 2s 29ms/step - loss: 0.5303 - accuracy: 0.7404
Final Loss: 0.5302808880805969, Final Accuracy: 0.7404426336288452

```

▼ Prediction

```

[ ] y_pred = model.predict(X_test, batch_size=batch_size)

```

```

249/249 [=====] - 4s 9ms/step





```


▼ First 10 predictions

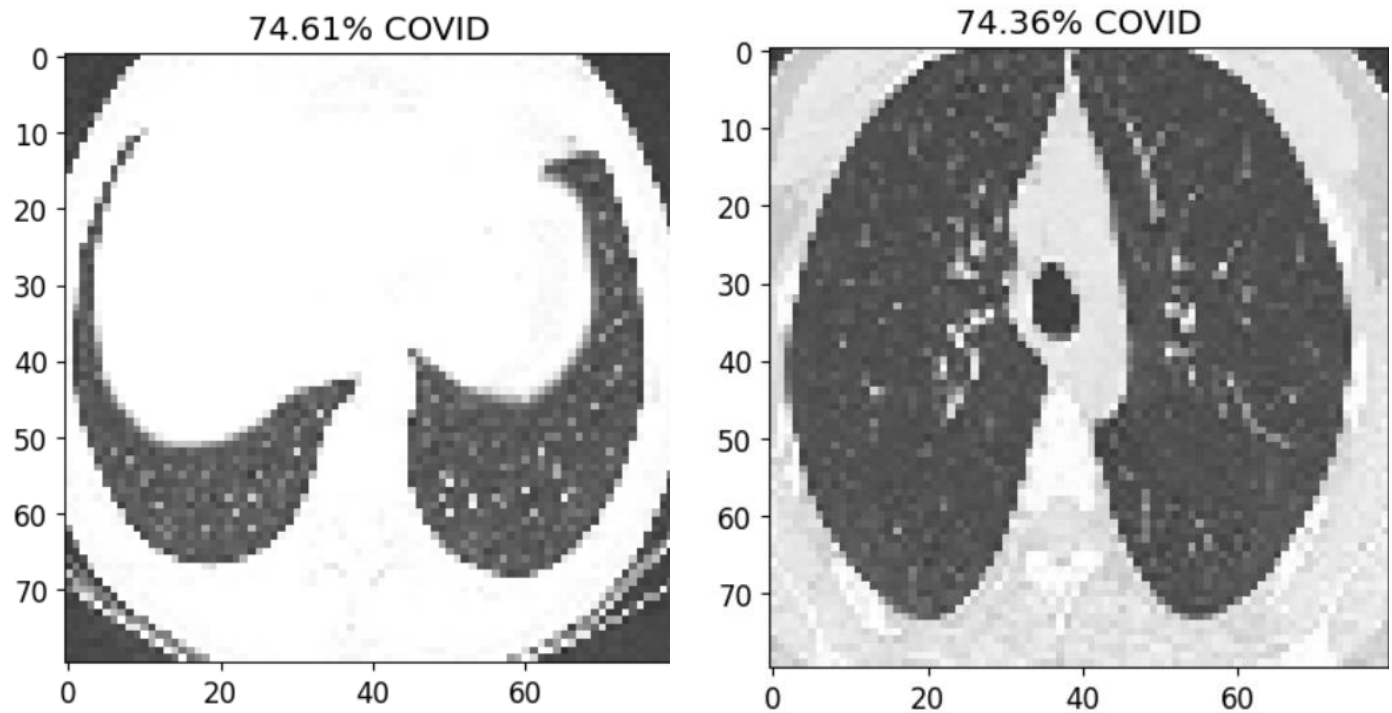
```

[ ] prediction=y_pred[0:10]
for index, probability in enumerate(prediction):
    if probability[1] > 0.5:
        plt.title('%0.2f' % (probability[1]*100) + '% COVID')
    else:
        plt.title('%0.2f' % ((1-probability[1])*100) + '% NonCOVID')
plt.imshow(X_test[index])
plt.show()

```





```
[ ] # Convert to Binary classes
    y_pred_bin = np.argmax(y_pred, axis=1)
    y_test_bin = np.argmax(y_test, axis=1)
```

▼ Classification Report

```
[ ] from sklearn.metrics import classification_report
    print(classification_report(y_test_bin,y_pred_bin))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.74 | 0.76 | 0.75 | 251 |
| 1 | 0.74 | 0.72 | 0.73 | 246 |
| accuracy | | | 0.74 | 497 |
| macro avg | 0.74 | 0.74 | 0.74 | 497 |
| weighted avg | 0.74 | 0.74 | 0.74 | 497 |


▼ Confusion Matrix

```
[ ] def plot_confusion_matrix(normalize):
    classes = ['COVID', 'NonCOVID']
    tick_marks = [0.5, 1.5]
    cn = confusion_matrix(y_test_bin, y_pred_bin, normalize=normalize)
    sns.heatmap(cn, cmap='plasma', annot=True)
    plt.xticks(tick_marks, classes)
    plt.yticks(tick_marks, classes)
    plt.title('Confusion Matrix')
    plt.ylabel('True label')
    plt.xlabel('Predicted label')
    plt.show()


print('Confusion Matrix without Normalization')
plot_confusion_matrix(normalize=None)

print('Confusion Matrix with Normalized Values')
plot_confusion_matrix(normalize='true')
```

Confusion Matrix without Normalization




Confusion Matrix with Normalized Values



▼ ROC Curve

```
[ ] fpr, tpr, thresholds = roc_curve(y_test_bin, y_pred_bin)
plt.plot(fpr, tpr)
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.0])
plt.rcParams['font.size'] = 12
plt.title('ROC curve for our model')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.grid(True)
```




Model Accuracy Plot

```
[ ] plt.plot(hist.history['accuracy'])
    plt.plot(hist.history['val_accuracy'])

    plt.title('Model Accuracy')
    plt.ylabel('Accuracy')
    plt.xlabel('Epoch')

    plt.legend(['Training', 'Testing'])
    plt.savefig('resnet_ct_accuracy.png')
    plt.show()
```



Model Loss Plot

```
▶ plt.plot(hist.history['loss'])  
plt.plot(hist.history['val_loss'])  
  
plt.title('Model Loss')  
plt.ylabel('Loss')  
plt.xlabel('Epoch')  
  
plt.legend(['Training', 'Testing'])  
plt.savefig('resnet_ct_loss.png')  
plt.show()
```

